

# Relational translation-II (Relationships)

# Knowledge objectives

1. Enumerate the different possibilities to enforce integrity constraints in a database
2. Explain the problem generated on dropping a table from the logical schema, even if this does not contain other attributes than the PK
3. Enumerate two possibilities to replace a table in the logical schema, when this does not have other attributes than the PK
4. Enumerate the three criteria to choose the PK on fusing two tables linked by a 1-1 association
5. Explain the two possible implementations of symmetric reflexive associations in a RDBMS
6. Remember where to place the attributes of the UML associations when they are implemented on a RDBMS, depending on their multiplicity
7. Distinguish an FK with two attributes from two FKs of one attribute
8. Distinguish association classes that appear just due to UML syntax constraints from those truly association classes
9. Enumerate ten pros and cons of storing multivalued attributes either per row or per column

# Understanding objectives

1. Solve a deadlock in the definition of FKs
2. Solve a deadlock in the loading of tables with FKs
3. Translate from a UML class diagram (with around 10 classes, some maybe associative classes, and related by associations, generalizations and aggregations) into an SQL schema

# Application objectives

1. Choose and justify the best option to translate from a UML class diagram (with less than 10 classes, some maybe associative classes, related by associations, generalizations and aggregations) into an SQL schema, given the statistics of participation of the instances in the relationships and the queries
2. Given a multivalued attribute and an explanation of its usage, choose and justify the best option to implement it in a RDBMS

# Constraints

Candidate keys

Foreign keys

# Implementing constraints


1. In the CREATE TABLE sentence  
Usually available, efficient, automatic and internal
2. Assertions  
Efficient, automatic and internal
3. Persistent Stored Modules
  1. Triggers  
Automatic and internal
  2. Procedures/functions  
Internal
4. Call Level Interface (e.g., ODBC, JDBC)  
Always available

LOGICAL DESIGN  
PHYSICAL DESIGN

# Candidate keys

- Primary
  - May not be available in our DBMS (really rare)
  - Physically, generate a B-tree index
- Alternative
  - Are not part of standard SQL
  - Can be implemented by NOT NULL + UNIQUE
    - Physically, generate a B-tree index

# Deadlock in the definition of FK



Dept(dpt,...,head)  
 Empl(dni,...,dpt)

## A. Modify the tables

- 1) Create "Dept" without FK
- 2) Create "Empl" with FK
- 3) Alter "Dept" adding the FK


## B. Create three tables

- 1) Create "Dept" (it doesn't have FK)
- 2) Create "Empl" with its FK
- 3) Create "Head" with two FK

Dept(dpt,...)  
 Head(dpt,empl)  
 Empl(id,...,dpt)



# Deadlock in the load of FK



Dept( dpt,..., head)  
CS      1  
Empl( dni,..., dpt)  
1      CS

## A. Drop/Disable FK

- 1) Alter "Dept" to drop/disable FK
- 2) Insert the department
- 3) Insert all employees
- 4) Alter "Dept" to add/enable FK

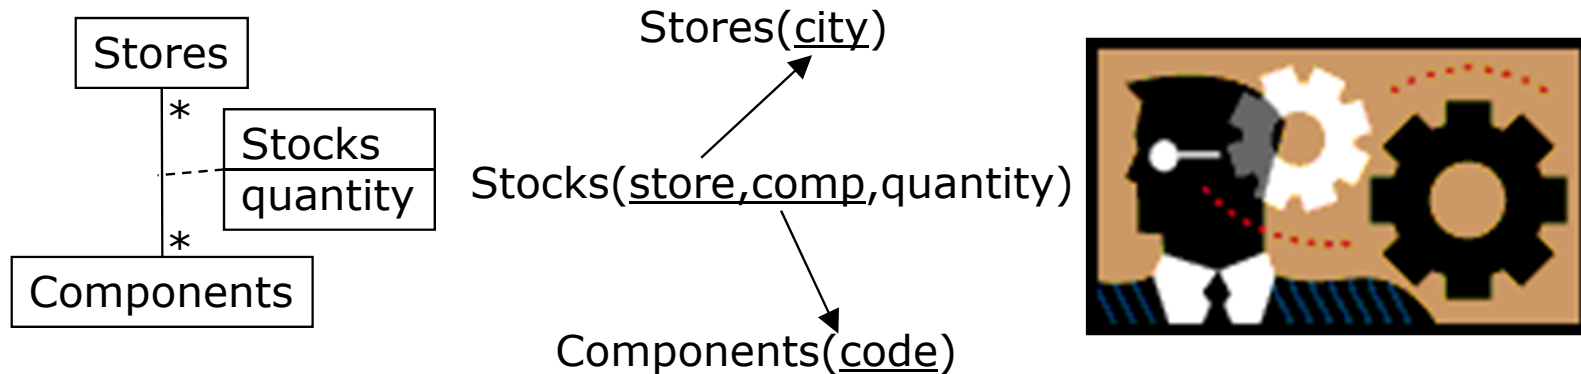
SET CONSTRAINT <name> [ENABLE|DISABLE];

## B. Defer the FK checking

SET CONSTRAINTS <name> [IMMEDIATE|DEFERRED];

# Relavance of FK

**Let's suppose that none of both classes has attributes**



Even if tables only contain the identifiers,  
they are still useful to enforce integrity constraints

**A table with few instances can be replaced by an enumerated data type or check constraint**

# Binary associations

# Multiplicities



## Maximum multiplicity:

- Question:
  - Each instance in one end, how many can have at most in the other end?
- Possibilities for binary associations:
  - $*_*$
  - $1_*$
  - $1-1$

## Minimum multiplicity:

- The three cases above split into subcases
- Question:
  - Could zeros exist (possible no participation of an instance in the association)?
    - If there are zeros, do they give rise to nulls in the relational translation?



# Binary association (\*-\*)



Projects(pro, ...)  
Assig(empl, pro, func)  
Employee(id, ...)



Always a new table!!!

# Binary associations (1-\*)



Dept(dpt, ...) **TRIGGER**

Empl(id, ..., dpt) **WITHOUT**





Dept(dpt, ...) **TRIGGER**

Empl(id, ..., dpt) **WITH**

Dept(dpt, ...) **TRIGGER**



Belong(dpt, empl)

Empl(id, ...) **WITHOUT**



Dept(dpt, ...)

Empl(id, ..., dpt) **WITHOUT**



Dept(dpt, ...)

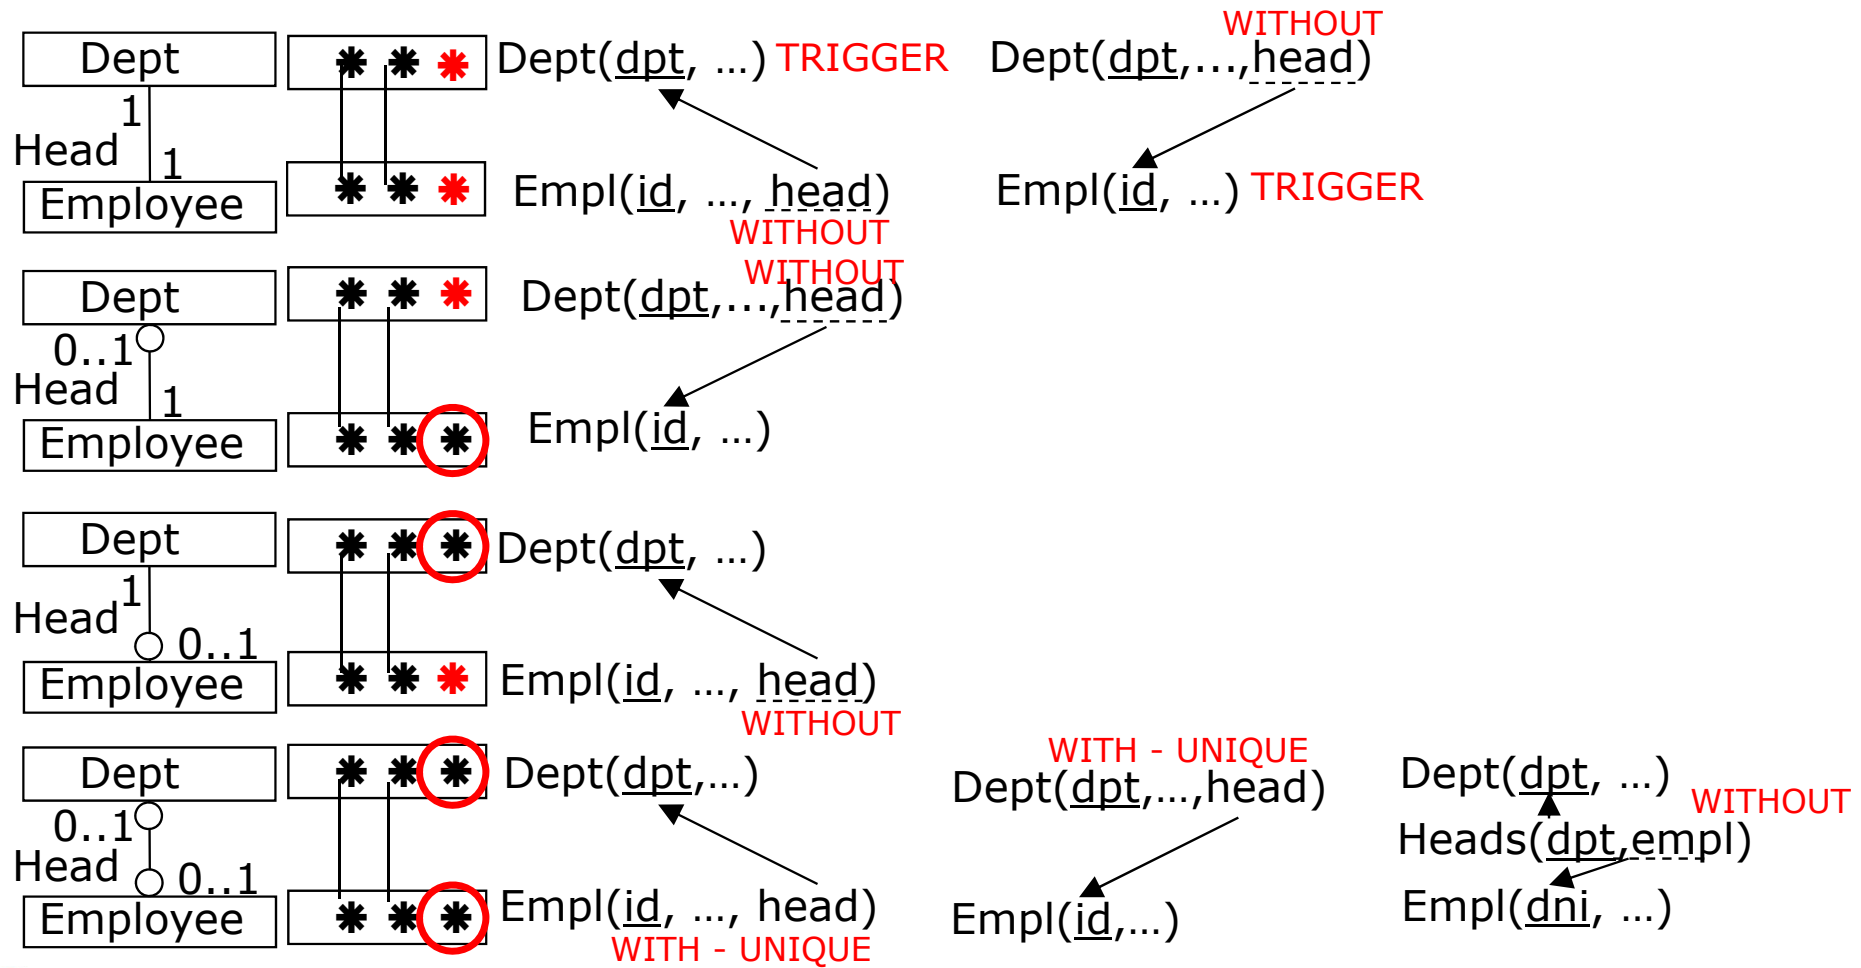
Empl(id, ..., dpt) **WITH**

Dept(dpt, ...)


Belong(dpt, empl)

Empl(id, ...) **WITHOUT**

# Binary associations (1-1)



# Fusing classes/Choosing PK



We have to choose the primary key among the candidate keys

The choice depends on:


- Frequency of use in the queries
- Space required
- Frequency of change

Country_President(	country, ...,	president, ...)
	USA	B. Obama
	Spain	M. Rajoy



# Reflexive associations

# Reflexive associations




- Valid multiplicities:

- $*-*$  (Relatives)
- $1-*$  (Mother)
- $1-1$  (Couple)

- Singularities:

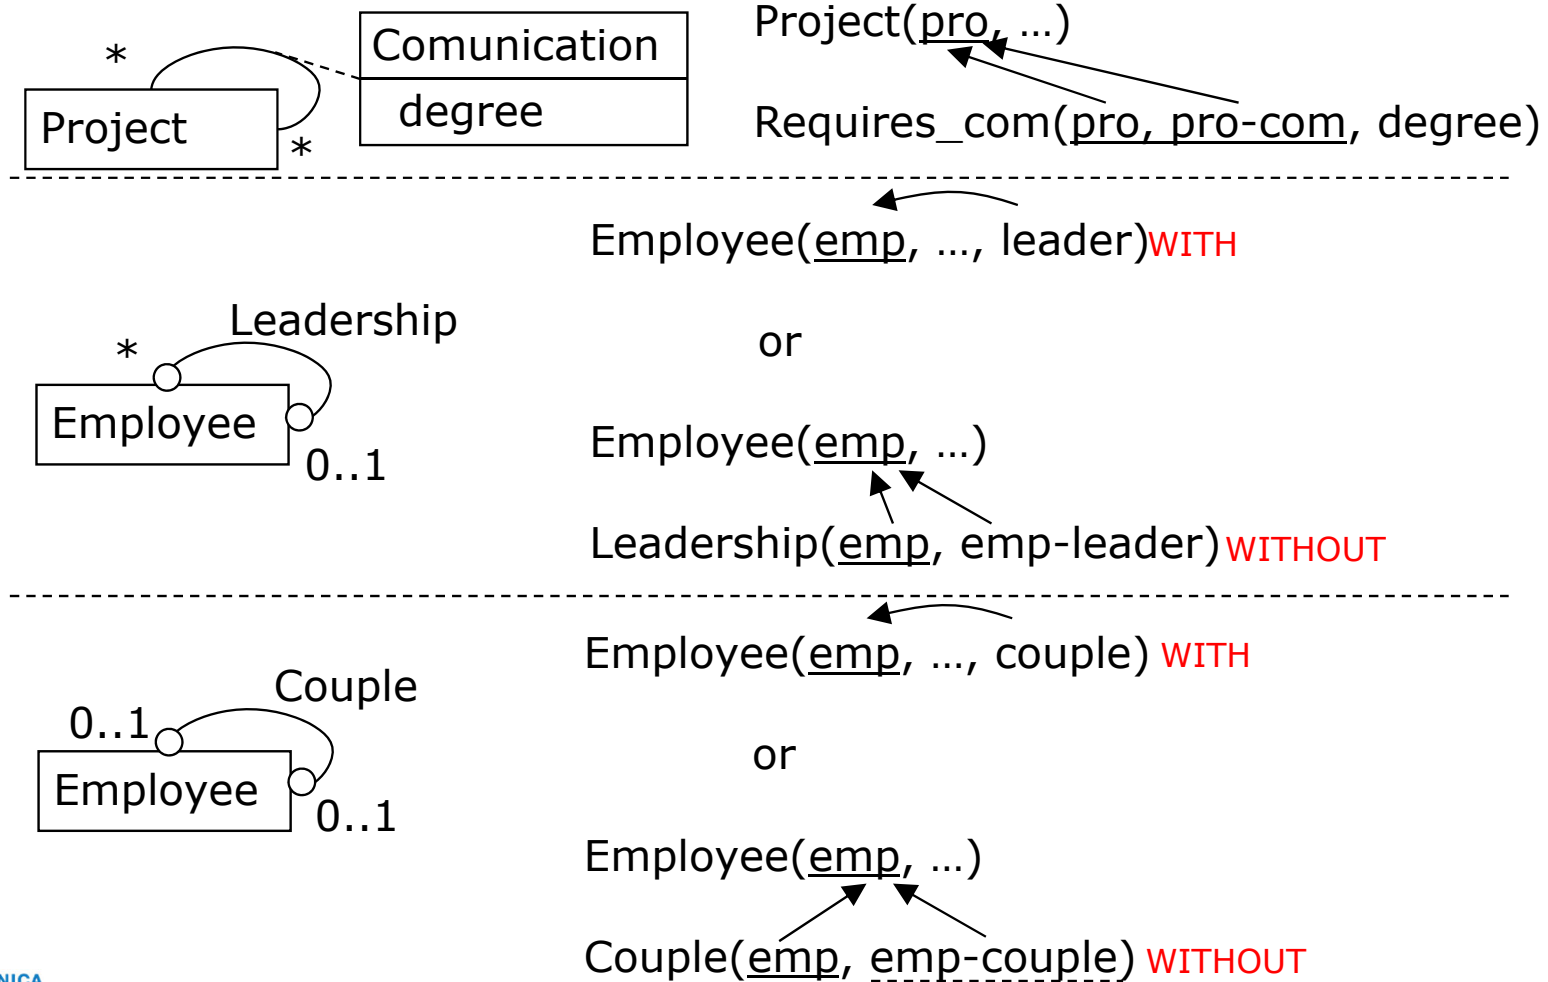
- Almost always have zeros
- May be symmetric or not



Brother1	Brother2
John	Peter
Peter	John

Friend1	Friend2	Grade
John	Peter	10
Peter	John	2

# Reflexive multiplicities



# Symmetric reflexive associations

- We must enforce the property
- Triggers may bring a satisfactory solution:

## a) Storing both pairs

ON INSERT (a, b) → INSERT (b, a)  
ON DELETE (a, b) → DELETE (b, a)  
ON UPDATE...

## a) Storing only one of the pairs

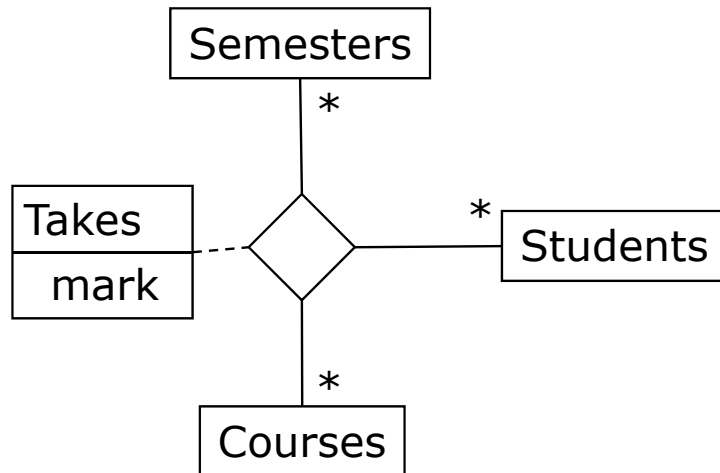
Simulate the whole set of pairs with a view:

```
CREATE VIEW v_couples AS  
SELECT emp1, emp2 FROM couples  
UNION  
SELECT emp2, emp1 FROM couples;
```

ON INSERT (a, b) → if (b, a) already present, raise exception  
ON DELETE (a, b) → if (b, a) is present instead, delete that  
ON UPDATE...

# N-ary associations


# Ternary associations (\*-\*-\* )



Students(st, ...)  
Semesters(sem, ...)  
Courses(cou, ...)  
Takes( st, sem, cou, mark)

Always a new table!!!

# Ternary associations (\*-\*1)



Parties(acronym, ...)

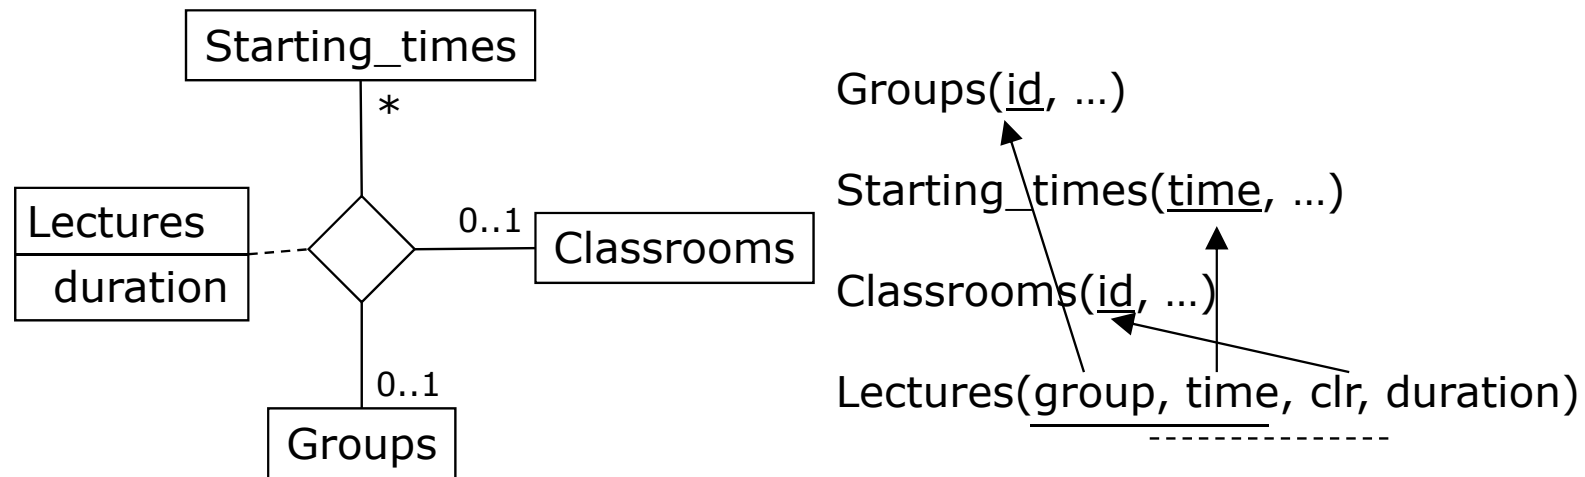
Aff\_dates(date, ...)

Politicians(name, ...)

Affiliates( par, date, pol, dis)

Always a new table!!!

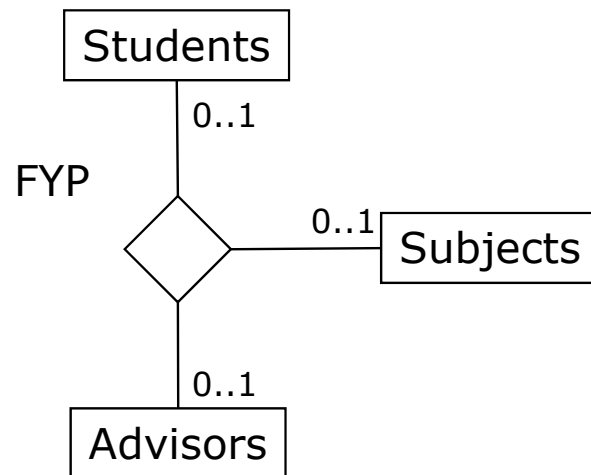
# Ternary associations (\*-1-1)



Always a new table!!!



# Ternary associations (1-1-1)



Subjects(subjects, ...)  
Students(name, ...)  
Advisors(name, ...)  
FYP( subj, student, advisor)  
-----

Always a new table!!!


# N-ary associations

- Binary: A new table or foreign key
- Ternary: A new table
- Quaternary: A new table
- ...


# Compound aggregation

# Compound aggregation (I)

- Weak class, with regard to the external key of the classical relational model




- A given class cannot be part of two



- Composition cannot have zeros at "to-one" side

# Compound aggregation (II)

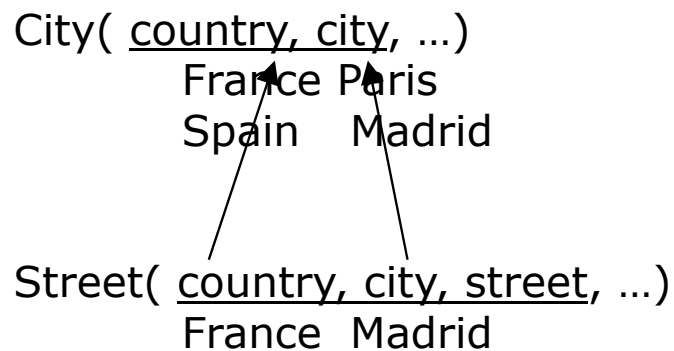


Country(coName, ...)  
City(coName, ciName, ...)  
Street( country, city, street, ...)

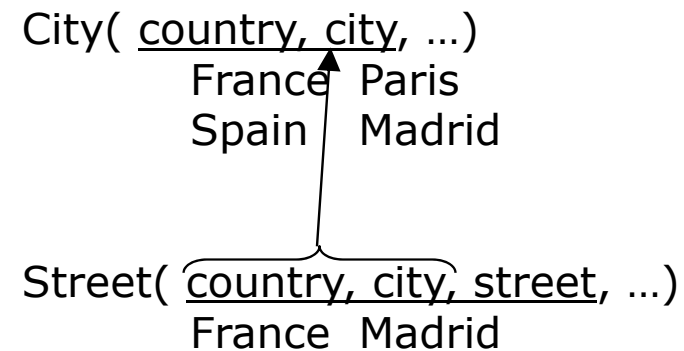
MULTIATTRIBUTE FK

# Multiattribute foreign keys

## Accepted



## FK violation




# Association classes




# Association vs class

- Association (common)
  - Identified through its “legs”
  - Makes defining/checking constraints difficult
- Class (uncommon)
  - Identified by itself (it has an external key)
  - Makes defining/checking constraints easy

## Association



## Class





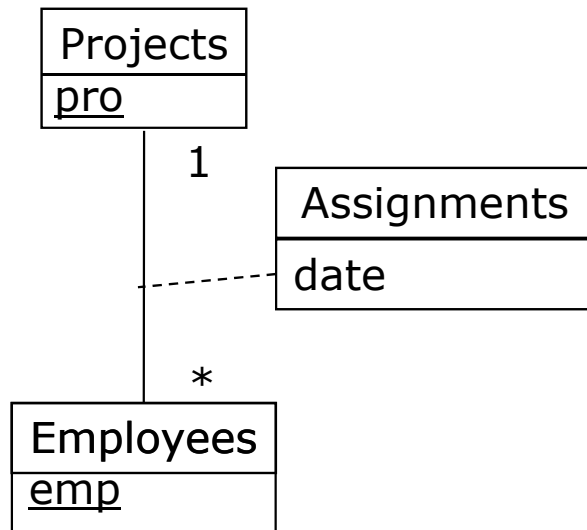
# Attributes of associations

- $*-*$  or n-ary (common)
  - In the table representing the association
- $1-*$  (uncommon)
  - If any, in the table representing the association
  - Otherwise, in the table at the side of the star
- $1-1$  (rare)
  - If any, in the table representing the association
  - If only one table (fusion), in it
  - Otherwise, in the table at the side of the zero (if any)

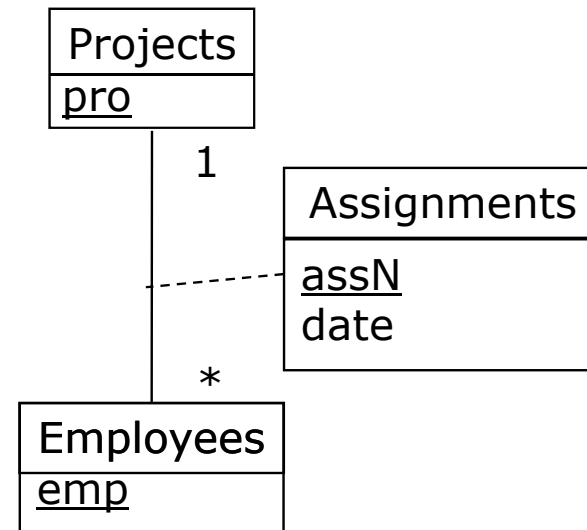
**Always together with the foreign key**

# Association vs association class

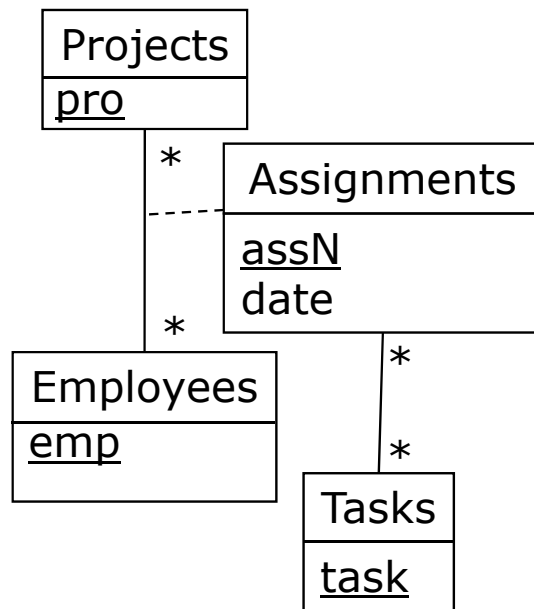
## Association



## Association class



# Association classes (I)




Implemented as a table

Projects( pro, ... )  
Employees( emp, ... )  
Assignments( assN, date, emp, proj, date )  
Task-Assig( task, assN )  
Task( task, ... )

Arrows indicate foreign key relationships: from assN in Assignments to pro in Projects, from assN in Assignments to emp in Employees, from assN in Task-Assig to assN in Assignments, and from task in Task-Assig to task in Task.

FOREIGN KEY (assN) REFERENCES Assignments (assN)

# Association classes (II)



The FK cannot point to "emp", because a task must be done by an employee who is assigned


Implemented as an FK

Projects(proj, ...)  
Employees(emp,... , proj, assN, date)  
Task-Assig(task, assN)  
Tasks(task, ...)

UNIQUE(assN)

CHECK ((proj IS NULL AND assN IS NULL) OR (proj IS NOT NULL AND assN IS NOT NULL))

# Association classes (III)



The FK cannot point to "Employees", because a task must be done by an employee who is assigned

Implemented as a table

Projects(proj, ...)  
Employees(emp,... )  
Assignment(proj, emp, assN, date)  
Task-Assig( task, emp )  
Tasks(task, ...)

proj NOT NULL

# Multivalued attributes

Per column

Per row



# Multivalued attributes example (I)

Client
<u>code</u> : integer phone: string [*] ...

One value per column

Client( code, office-phone, secretary-phone, cell-phone, ... )  
C1 933333333 933333331 666666666 null

One value per row

Client(code, ...)  
↑  
ClientTelephones( code, place, telephone )  
C1 Office 933333333  
C1 Secretary 933333331  
C1 Cellular 666666666

# Multivalued attributes example (II)

Experiment
<u>id</u> : integer measure: integer [*] ...

One value per column

Experiment(id, measure<sub>equipment1</sub>, measure<sub>equipment2</sub>, ..., measure<sub>equipmentn</sub>)  
a            30                            31                            28

One value per row

Experiment(id, equipment, measure)  
a        1        30  
a        2        31  
...  
a        n        28

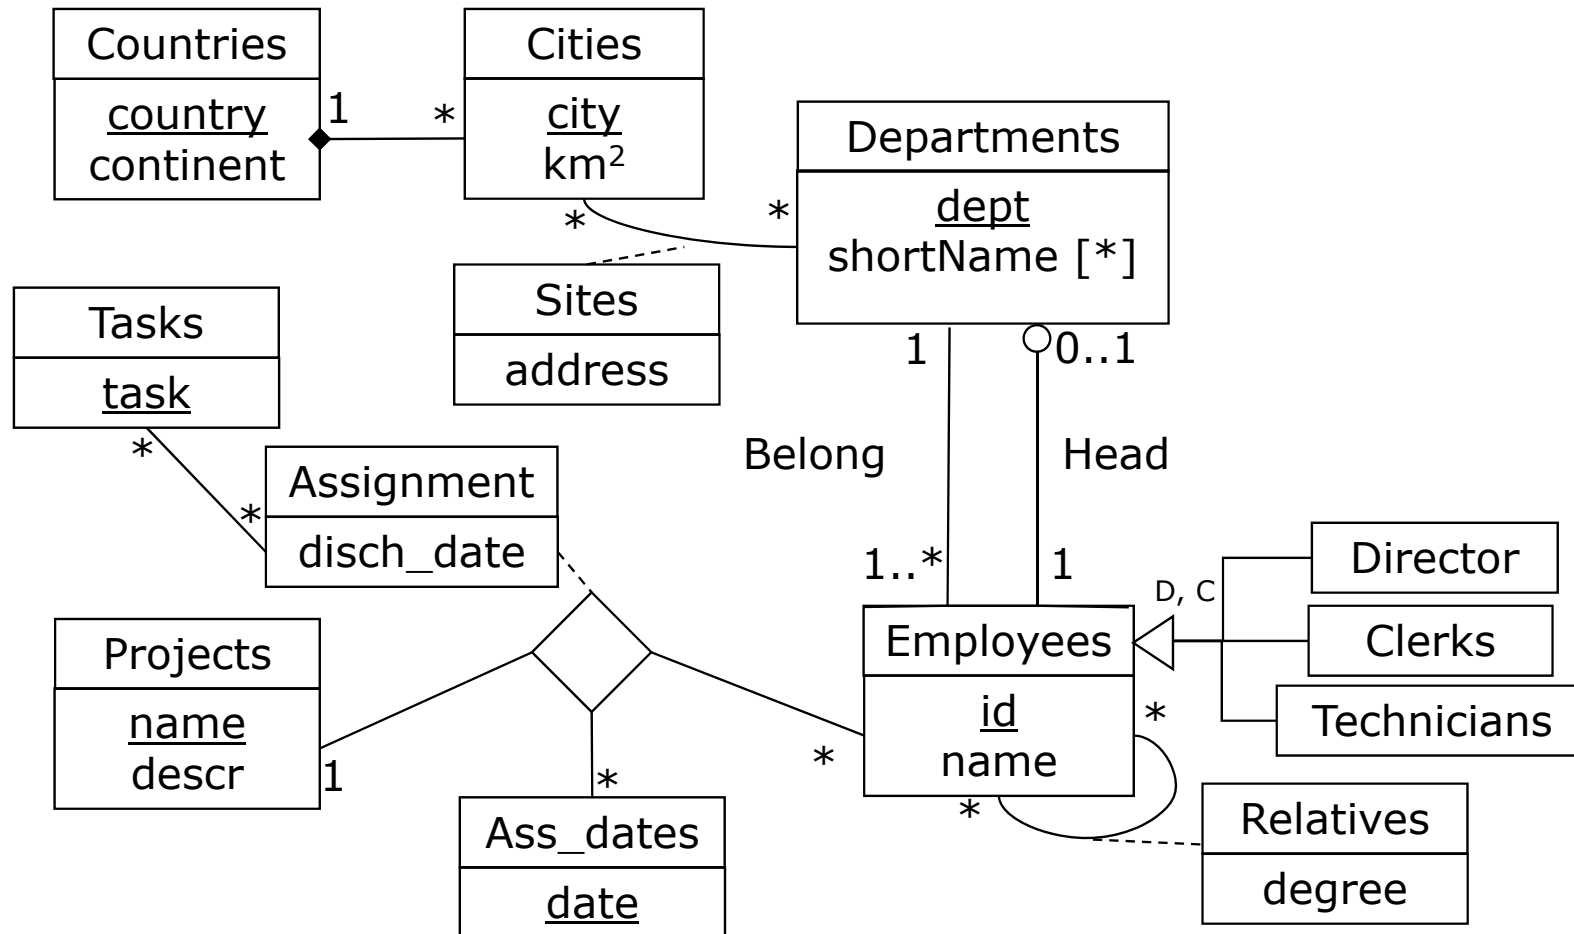


# Multivalued attributes comparison

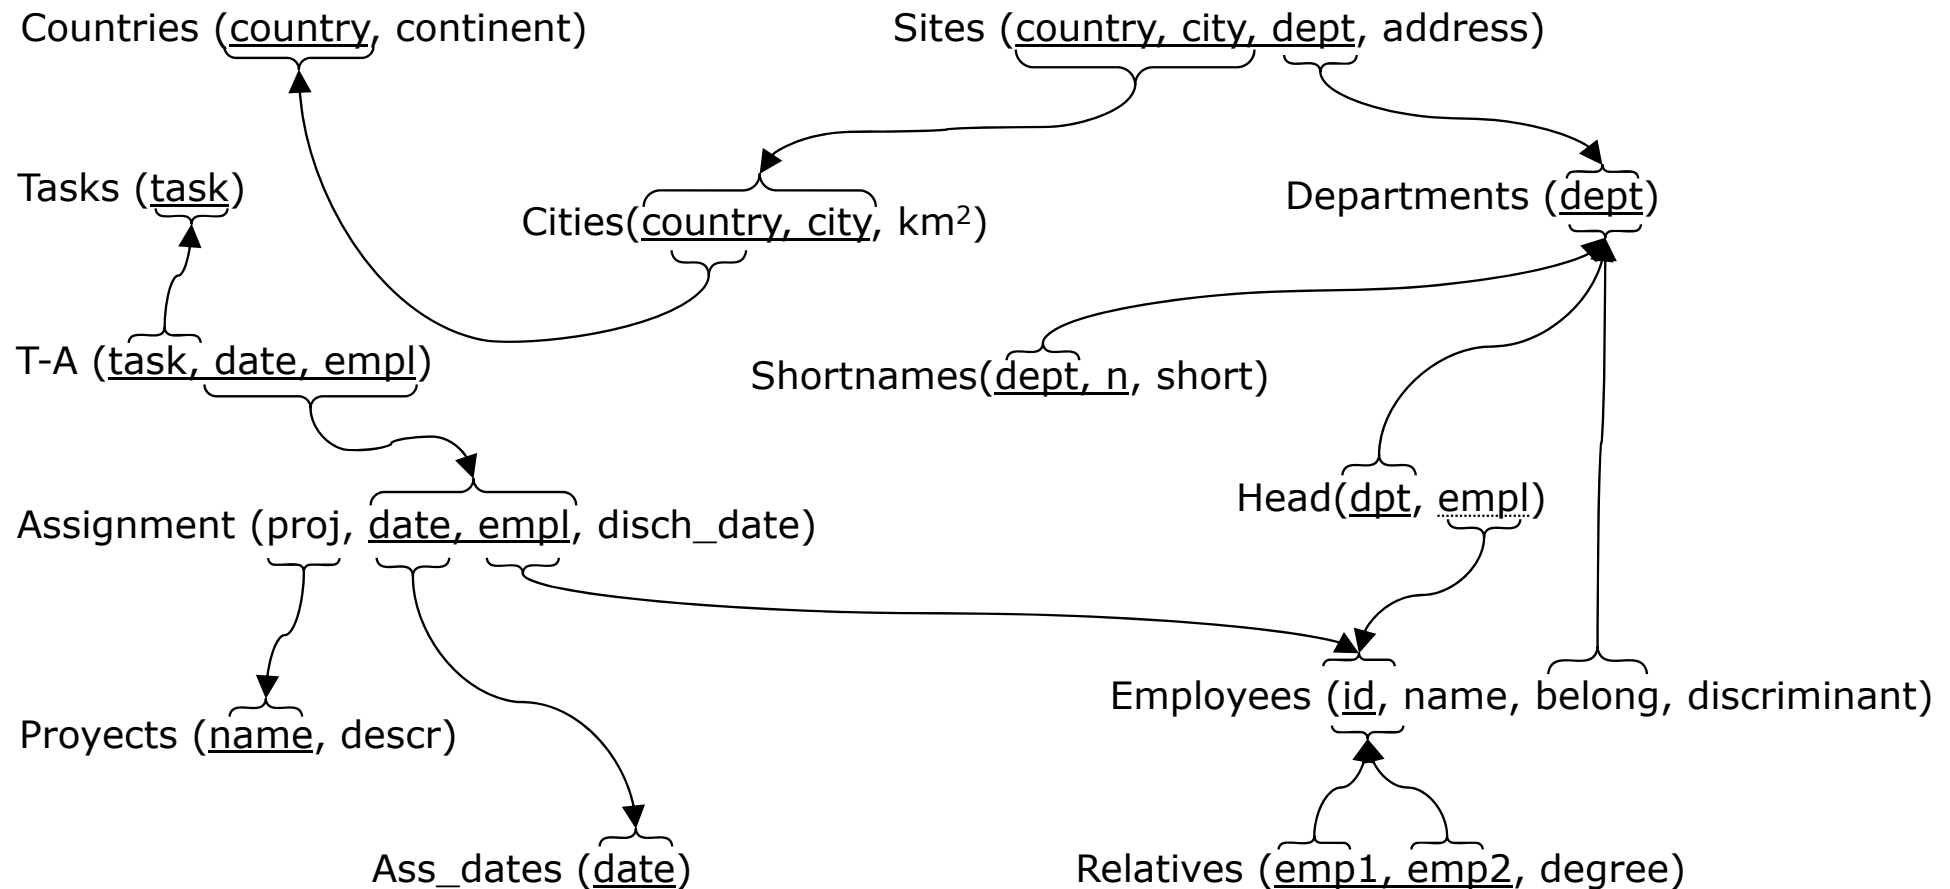
Per column	Per row
Fixed number of values	Variable number of values
Few values	Many values
Generates nulls	There are no null values
One I/O	Many I/O
Global processing	Partial processing
Natural PK	Artificial PK
Less space	More space
Hard to aggregate	Easy to aggregate
Many CHECKs	One CHECK
Lower concurrency	Higher concurrency

# Example of translation

# Example of conceptual schema




# Example of logical schema



# Example of constrain implementation

Constraints reflected in the conceptual schema:

1. An employee is head of zero or one department
  - Attribute "empl" in table "Head" is UNIQUE
2. Every department has exactly one head
  - Attribute "empl" in table "Head" is FK and NOT NULL
  - An assertion should be defined to check that each and every department has a head
    - We may also have implemented the association with two tables instead of three
3. An employee belongs to exactly one department
  - Attribute "belong" in table employee has been defined as FK and NOT NULL
4. Each and every department has at least one employee
  - An assertion should be defined



Dept(dept,...)  
Head(dept,empl)  
Empl(id,...,belong)

# Closing

# Summary

- Constraints
  - Deadlocks
  - Relevance of FK
- Translation of relationships
  - Binary
    - \*\_\*
    - 1-\*
    - 1-1
  - Reflexive
  - N-ary
  - Compound aggregation
  - Association classes
    - Attributes of relationships
- Multivalued attributes

# Bibliography

- J. Sistac. *Sistemes de Gestió de Bases de Dades*. Editorial UOC, 2002
- J. Sistac et al. *Disseny de bases de dades*. Col·lecció Manuals, number 43. Editorial UOC, 2002
- R. Elmasri and B. Nbathe. *Fundamentals of Database Systems*. Addison-Wesley, 4th edition, 2003
- T. Teorey et al. *Database modeling and design*, 4<sup>th</sup> edition. Morgan Kaufmann Publishers, 2006