

# Views

# Knowledge objectives

1. Explain the differences between the three levels in the ANSI/SPARC architecture, paying special attention to physical and logical independency
2. Explain the two differences between a view and a table
3. Explain the difference between a view and a materialized view
4. Name four potential uses of views
5. Enumerate and distinguish the four problems associated to views
6. According to the standard, name the two properties a view must fulfill to be updatable
7. Enumerate when and how a materialized view can be refreshed
8. Discuss the benefits of a complete and an incremental view update
9. Enumerate the requirements for a query to be rewritten in terms of a materialized view
10. Identify the two problems in pre-computing queries

# Understanding objectives

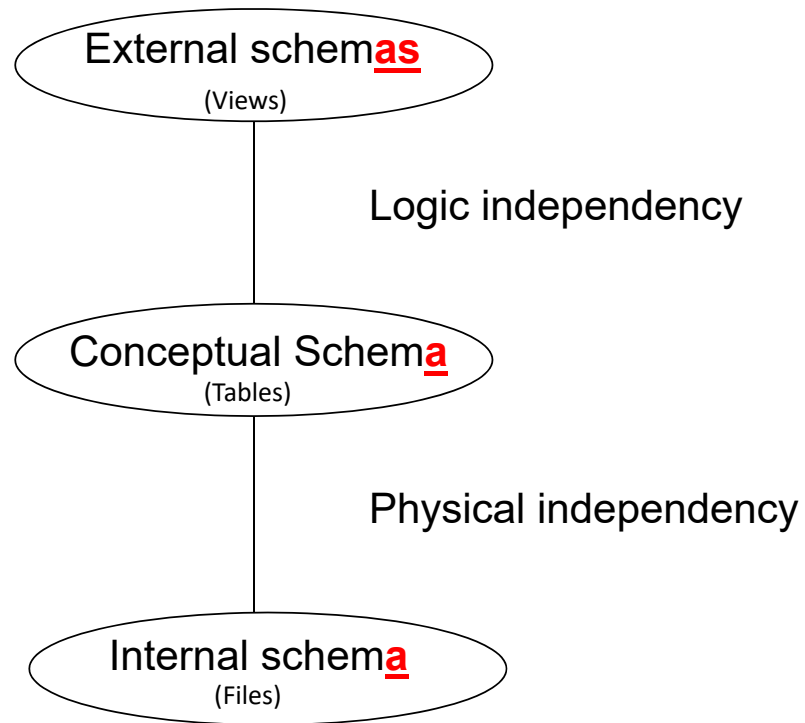
1. Given a set of tables and some views, apply view expansion over a query
2. Given the number of dimensions and levels in each of them, calculate the number of different materialized views in a multidimensional schema, considering only the cases in the group by clause
3. Select a set of views to be materialized using a greedy algorithm in the following scenarios:
  - a) Disk space is limited and the system is read-only
    - a) Only the given user queries can be materialized
    - b) Any query can be materialized
  - b) Disk space is not limited and the system is read-write

# Application objectives

1. Given an assertion simulate its implementation using the materialized view syntax in Oracle
2. Given a set of source tables (no more than 6) together with their statistics and some views over them (no more than 3), justify if
  - a) A given view is updatable
  - b) It is worth or not an incremental update of a materialized view in front of a complete one
  - c) A specific query over the tables can be rewritten in terms of the materialized views

# View definition and difficulties

# ANSI/SPARC architecture



# Alternatives to implement a relation

- From the structures / data point of view
  - Tables
    - Data in disk (i.e., Materialized)
  - [Non-materialized] views
    - Definition in catalog (SQL statement)
      - Re-executed with every query
  - Materialized views
    - Data in disk (i.e., Materialized) and definition in catalog (SQL statement)
- From data retrieval point of view
  - Tables
    - Querying the materialized data
  - [Non-materialized] views
    - Transforming the query into another one over the underlying tables
$$V = F(R_1, R_2, \dots, R_n)$$
$$Q(V) \rightarrow Q'(R_1, R_2, \dots, R_n)$$
  - Materialized views
    - Querying the pre-computed (i.e., Materialized) result of the query
      - Reduced to a synchronization problem

# Example in Data Warehousing

## Dimension

STORES	
<u>Id</u>	city
1	Mataró
2	Mataró

## Fact

SALES				
<u>storeId</u>	<u>date</u>	<u>time</u>	<u>productId</u>	euros
1	xxx	xxx	1	10
2	xxx	xxx	1	15
1	xxx	xxx	2	20

## Dimension

PRODUCTS	
<u>Id</u>	product
1	Rubber
2	Pen

## Materialized view

EUROSALES			
<u>City</u>	<u>Product</u>	sumEuros	salesCounter
Mataró	Rubber	25	2
Mataró	Pen	20	1

Measure

Aggregations



# Potential uses


- Simplify complex schemas
  - Simplify queries
- Hide data / implementation details
  - Solve security issues
- Improve performance
  - Only if materialized
- Integrity checking

# DBMS difficulties/features associated to views

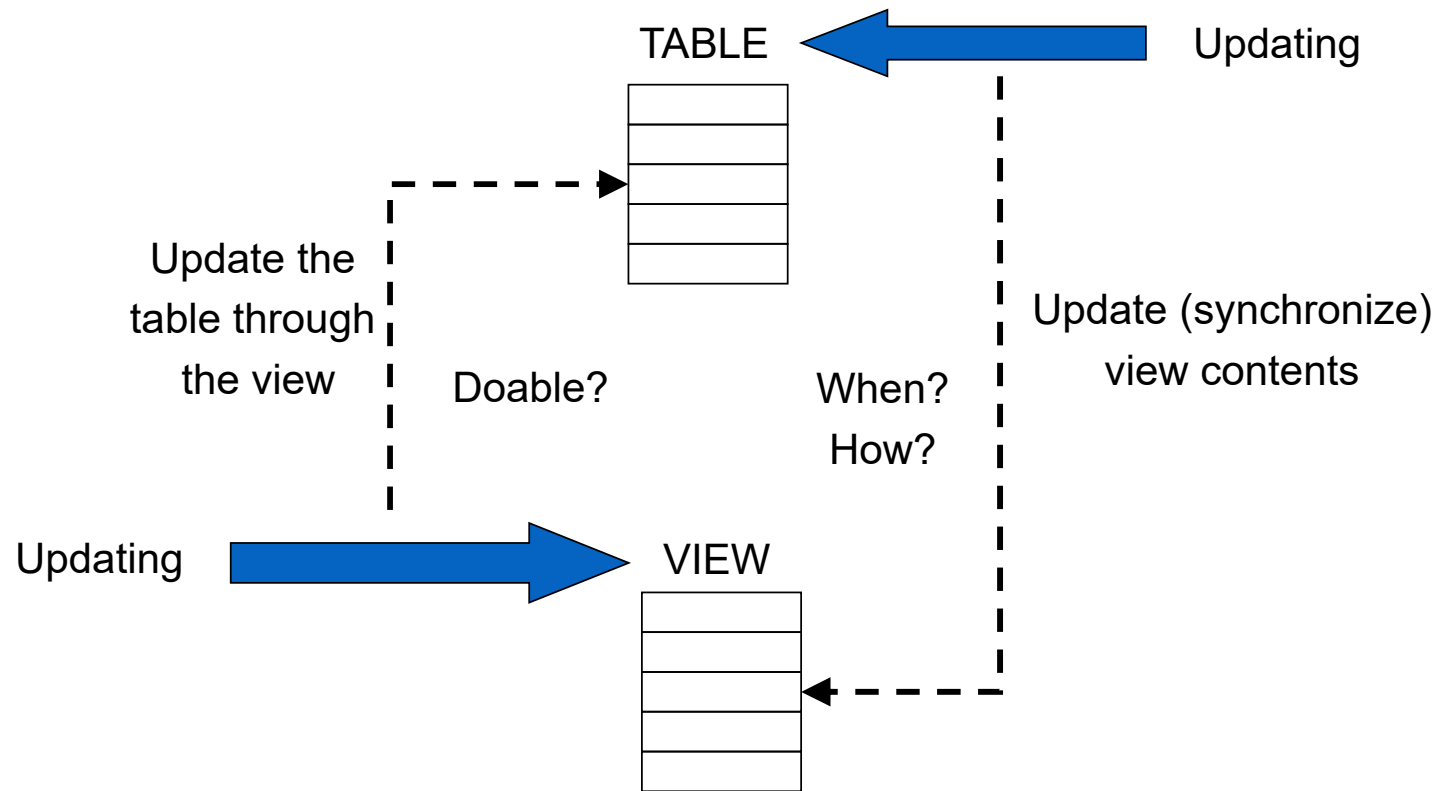
	Queries	Modifications
Over a View → Affect a Table	View expansion	Update through views
Over a Table → Affect a (Materialized) View	Query rewriting	View updating

- Non materialized views
  - a) View expansion
    - Transform the query over the views into a query over the source tables
- Materialized views
  - b) Query rewriting (i.e., answering queries using views)
    - Transform an arbitrary query over the tables into a query over the available views
  - c) View updating
    - Changes in the sources are, potentially, propagated to the view
- Both
  - d) Update through views
    - Propagate the changes in the view to the sources by means of a translation process

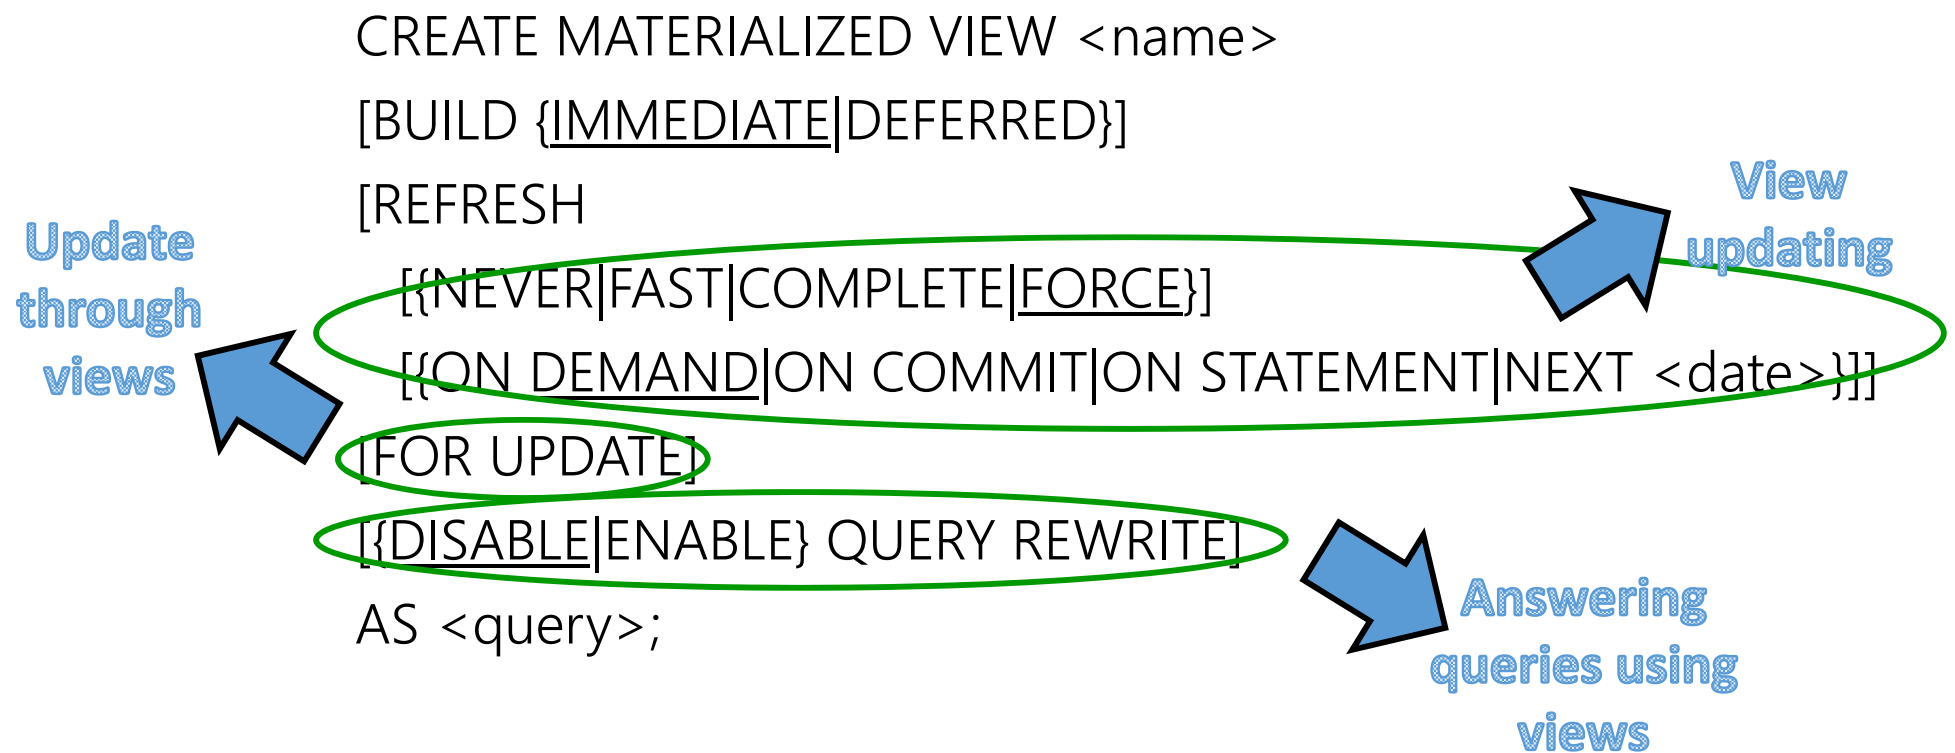
# Views and queries



# Views and modifications



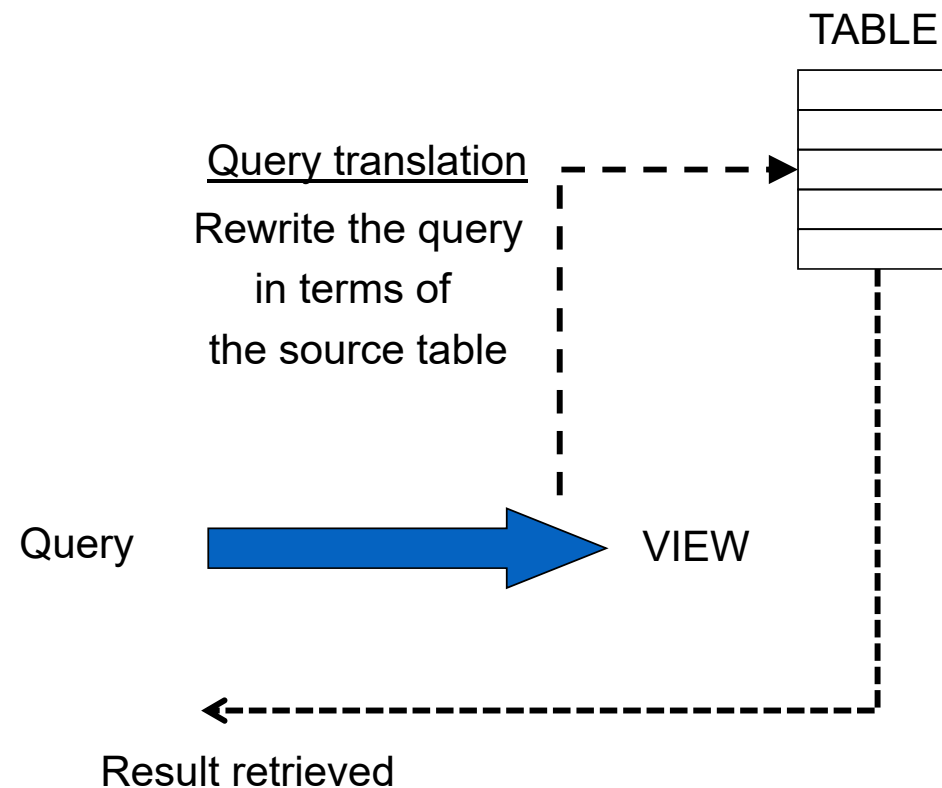
# Materialized views in Oracle (not in SQL:2023)



# View expansion

# View expansion problem

In case the query is over a view, a preliminary step to unfold the view definition is required (!)



# View expansion solution

Do

1. Retrieve the view definition(s) from the catalog
2. Replace the view definition(s) in the query

While unresolved views in the query



# Example of view expansion

```
CREATE VIEW richempl AS  
  SELECT *  
  FROM employees  
  WHERE salary > 30000;
```

```
SELECT AVG(salary)  
FROM richempl;
```

```
CREATE TABLE employees (  
  dni CHAR(8),  
  name CHAR(8),  
  salary INTEGER,  
  PRIMARY KEY (dni)  
);
```

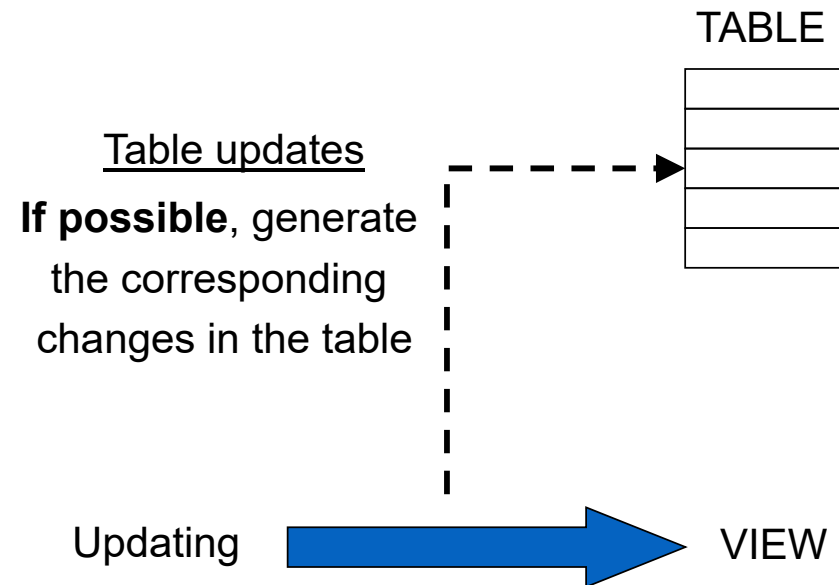


```
SELECT AVG(salary)  
FROM (  
  SELECT *  
  FROM employees  
  WHERE salary > 30000  
);
```

# Update through views

# Update through views problem

In case the update is over a view, it is still possible to modify the table to produce the desired effect, under some conditions (!)

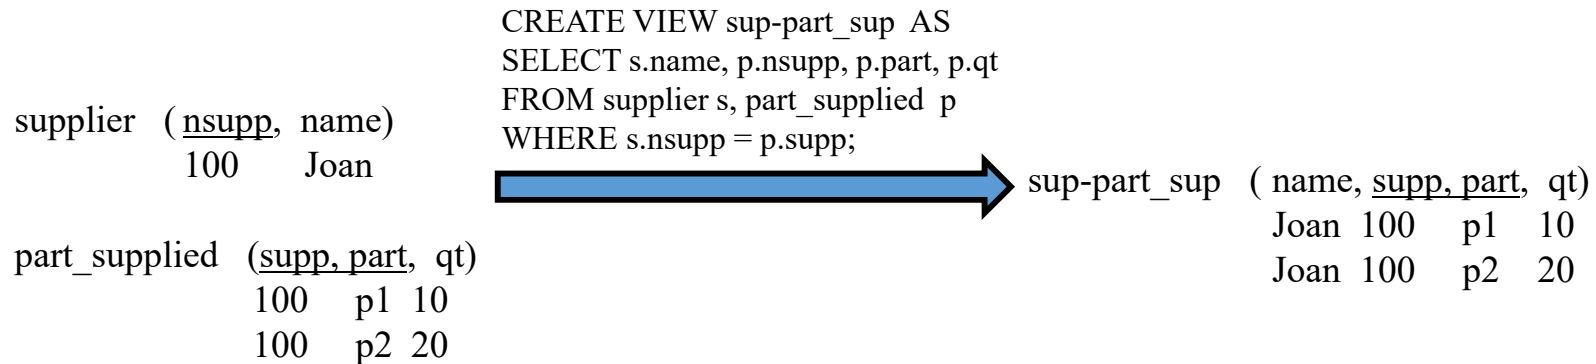


# Update through views solution

- Views, in general, are non-updatable
  - Only when the update can be translated unambiguously over the relational table
- According to the standard, views can be updated only if
  - a) It is an algebraic selection on a single relation or updatable view
  - b) It may also contain an algebraic projection iff the following are projected:
    - The primary key
    - All not-null attributes

Subqueries, joins or aggregate functions are not allowed!!!

# Example of update through views with joins



INSERT INTO sup-part\_sup VALUES ('Pere',200,p1,20);



INSERT INTO supplier VALUES (200,'Pere');  
INSERT INTO part\_supplied VALUES (200,'p1',20);


UPDATE sup-part\_sup SET name='Joana'  
WHERE supp=100 AND part='p1';



DELETE FROM sup-part\_sup WHERE supp=100 AND part='p1';



# Example of update through views with aggregates

part_supplied ( <u>supp</u> , <u>part</u> , qt)	CREATE VIEW total_qt (part, total) AS SELECT part, sum(qt) FROM part_supplied GROUP BY part;	total_qt ( <u>part</u> , qt)
sp1 p1 100		p1 300
sp2 p1 200		p2 200
sp2 p2 200		

INSERT INTO total\_qt VALUES ('p3',400);

UPDATE total\_qt SET qt=301 WHERE part='p1';

DELETE FROM total\_qt WHERE part='p1';

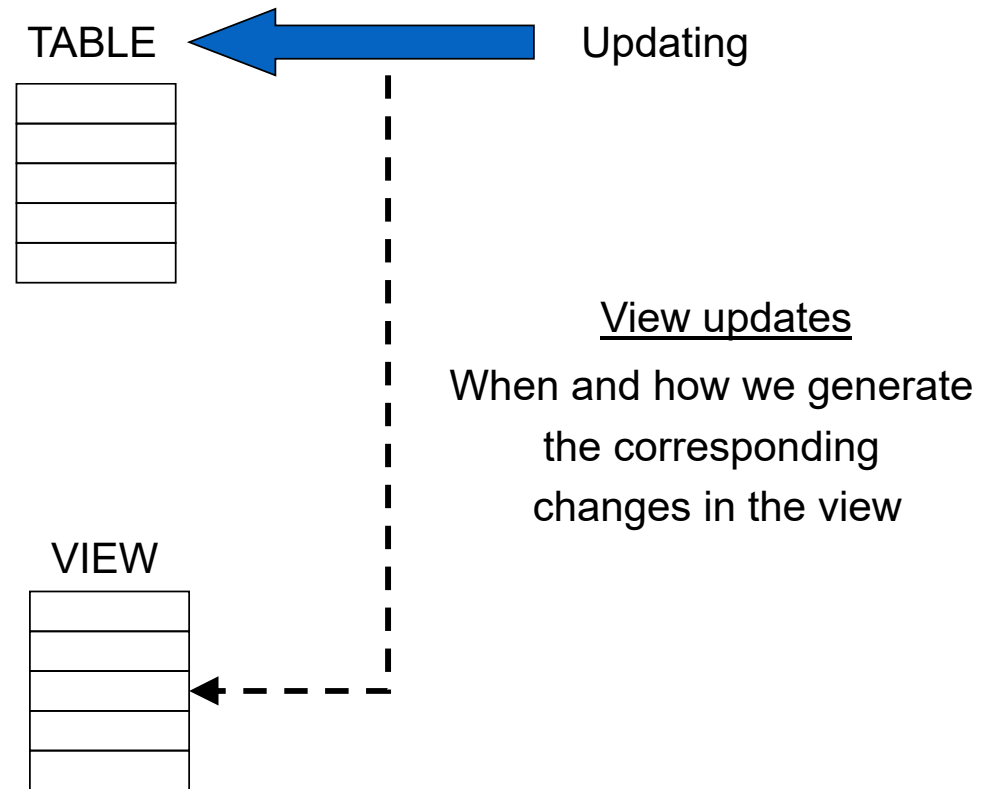


DELETE FROM part\_supplied  
WHERE part='p1';

# View updating

# View updating problem

In case the update is over a table, we should eventually synchronize the content of the view to avoid inconsistencies (!)





# View updating (when)

- On statement
  - Avoids the need of the associated log, but is very inefficient, in general
    - a) All DML operations are slower, as they include the update of the MV
      - The same MV row can be updated several times in the same transaction
    - b) The transaction cancellation is more expensive, as it includes undoing changes in the MV
- On commit
  - May still be too expensive
- On demand
  - Generates temporal inconsistencies
- Next <date>
  - Generates temporal inconsistencies

# View updating (how)

- Complete update

- All instances are regenerated

`REFRESH MATERIALIZED VIEW <name>;`

- Clearly inefficient ?
    - Always possible

- Incremental update (called “fast” in Oracle)

- Only instances that changed are regenerated
    - Much more efficient ?
    - Not always possible
      - Depends on the information available

# Incremental materialized view refresh (in Oracle)

- Incremental (called fast) updates allow on commit refreshment (otherwise not allowed)
- A log must be defined for every source table
  - Only one log per table is allowed!
  - Stores rows describing changes from last refresh
  - Tuples should be univocally identified (ROWID or PK needed)

```
CREATE MATERIALIZED VIEW LOG ON table
    [WITH [PRIMARY KEY,] [ROWID,] [SEQUENCE] (list_of_attr)]
    [{INCLUDING | EXCLUDING} NEW values]
```

- Both the log and the view definition query (Q') must fulfill a set of constraints depending on the kind of query
  - Basic queries (without groupings nor joins)
  - Join queries
  - Grouping queries

# Assertions

- A constraint that may involve several tuples or tables
  - Introduced in SQL:1992
  - They are, generically, not yet implementable in most RDBMS (!)
- Can be simulated using materialized views in some DBMS (e.g., Oracle)
  - There are two alternative implementations:
    - a) Using an empty materialized view
      - 1) Define an MV with the negation of the desired assertion definition
      - 2) Define a dummy check, which should never be satisfied (i.e., an inconsistent IC)
    - b) Using a non-empty materialized view
      - 1) Define an MV with the desired assertion definition without the where/having condition
      - 2) Define a check corresponding to the negation of the where/having of the assertion
- An ON COMMIT refresh will be required, most of the times
  - ON DEMAND or NEXT may be enough in some cases

# Example of assertion simulation (empty MV)

- Assertion (standard syntax):

```
CREATE ASSERTION IC_debt (NOT EXISTS
  (SELECT c.#customer
   FROM customers c, orders o
   WHERE c.#customer = o.#customer
        AND c.type = 'regular' AND o.payment = 'pending'
   GROUP BY c.#customer
   HAVING SUM(o.quantity) >= 10000));
```

- MV simulating the assertion (Oracle syntax):

```
CREATE MATERIALIZED VIEW mv BUILD IMMEDIATE REFRESH FAST ON COMMIT AS
  SELECT 'x' AS X
  FROM customers c, orders o
  WHERE c.#customer = o.#customer
        AND c.type = 'regular' AND o.payment = 'pending'
  GROUP BY c.#customer
  HAVING SUM(o.quantity) >= 10000))

ALTER TABLE mv ADD CONSTRAINT mv_check CHECK (X is null);
```

# Example of assertion simulation (non-empty MV)

- Assertion (standard syntax):

```
CREATE ASSERTION IC_debt (NOT EXISTS
  (SELECT c.#customer
   FROM customers c, orders o
   WHERE c.#customer = o.#customer
        AND c.type = 'regular' AND o.payment = 'pending'
   GROUP BY c.#customer
   HAVING SUM(o.quantity) >= 10000));
```

- MV simulating the assertion (Oracle syntax):


```
CREATE MATERIALIZED VIEW mv BUILD IMMEDIATE REFRESH FAST ON COMMIT AS
  SELECT c.#customer AS id, SUM(o.quantity) as debt
  FROM customers c, orders o
  WHERE c.#customer = o.#customer
        AND c.type = 'regular' AND o.payment = 'pending'
  GROUP BY c.#customer;

ALTER TABLE mv ADD CONSTRAINT mv_check CHECK (debt < 10000);
```

# Query rewriting

# Query rewriting problem

In case a regular query over a table, which has materialized views already defined on it, we could find execution strategies that use the materialized view(s) instead of the table (!)

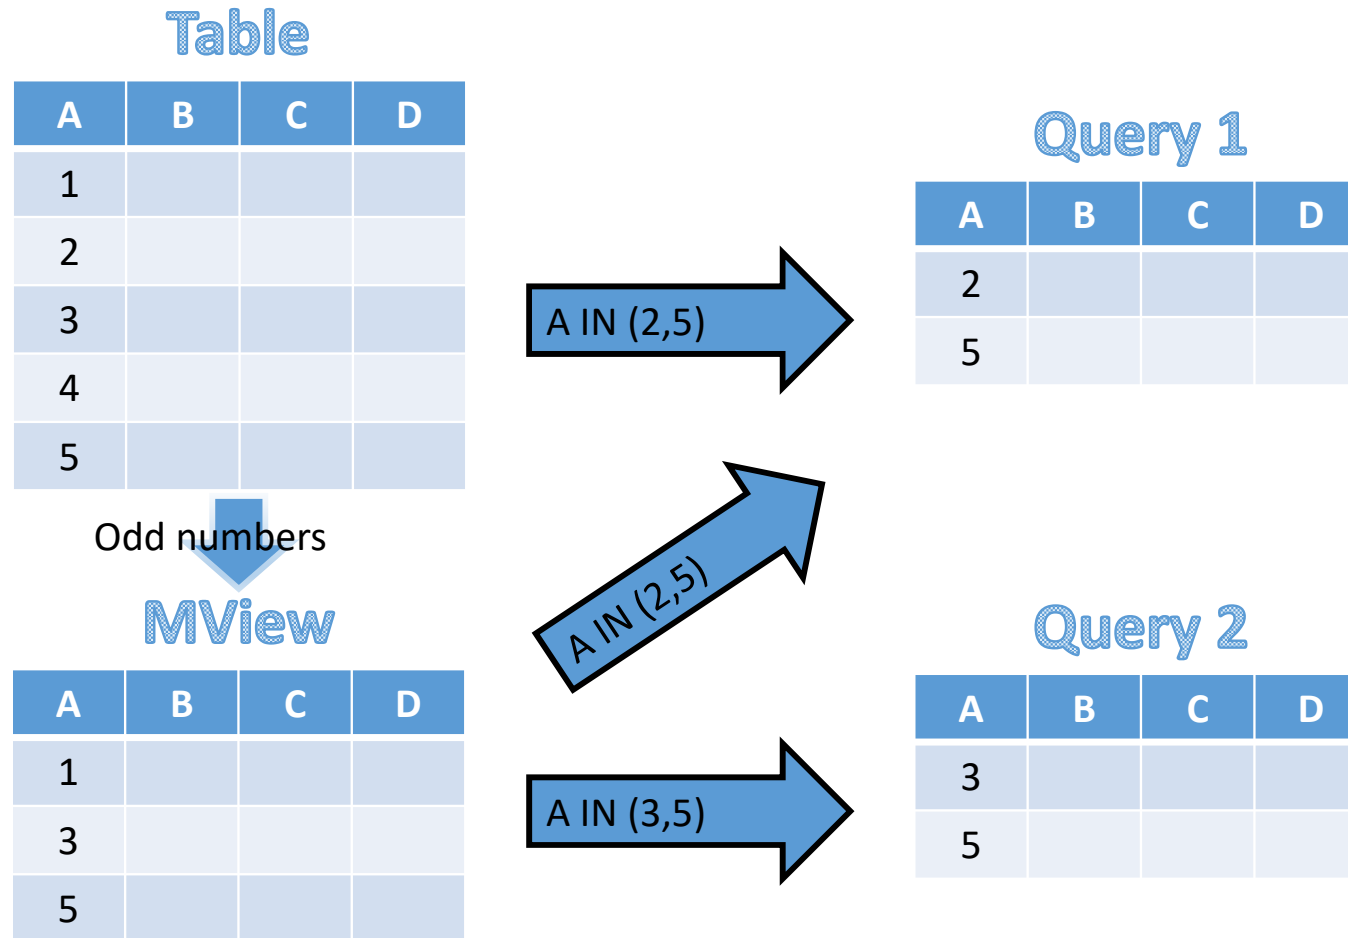




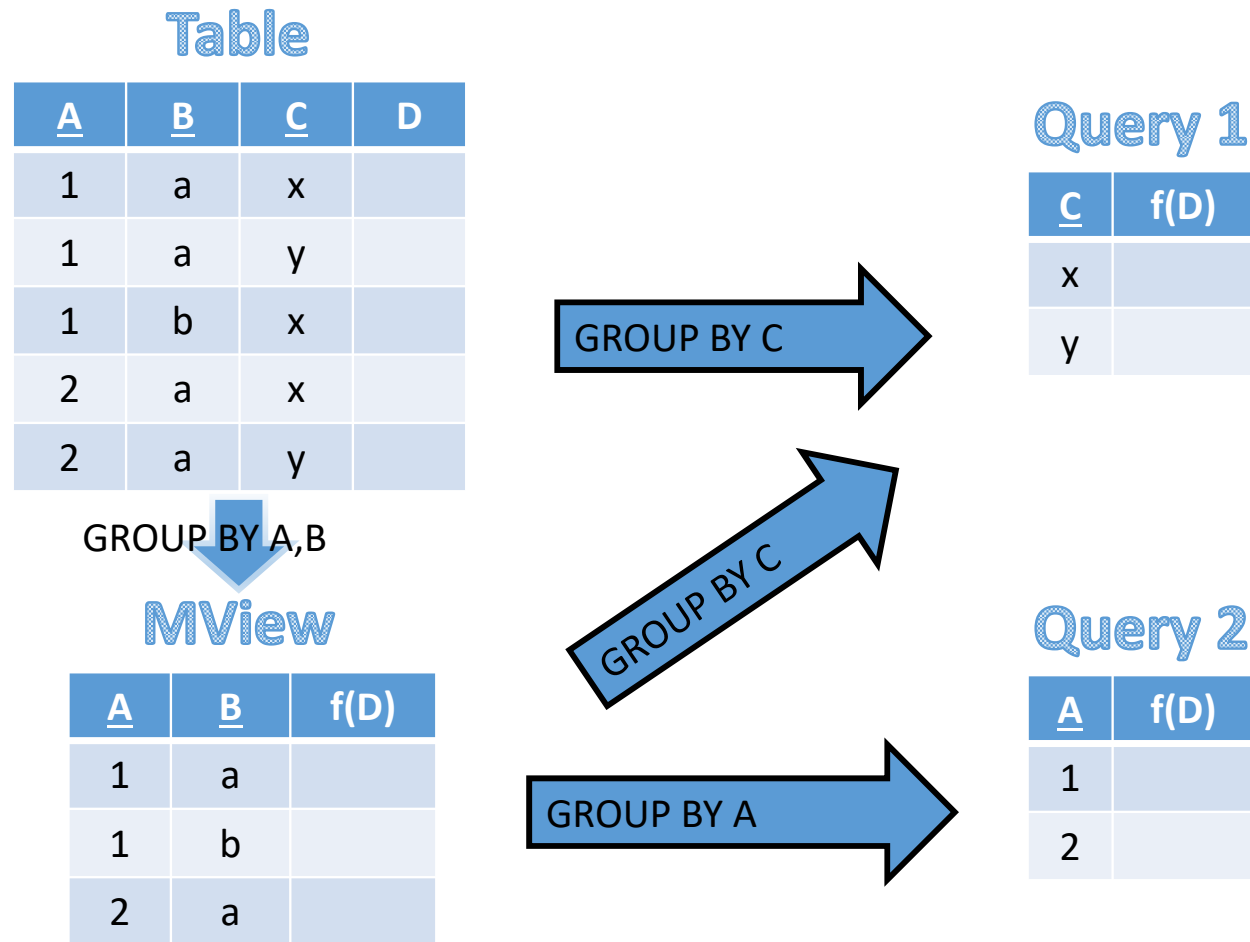
# Query rewriting solution

- Deciding whether it is possible to rewrite a query in terms of existing materialized views or not is computationally complex
  - Also known as answering queries using views
  - DBMSs restrict the search space to common cases by using rules
- Requirements:
  - a) Query predicate must be subsumed by that of the view  
 $\text{Query tuples} \subset \text{View tuples}$
  - b) Aggregation level in the query must be higher or equal to that in the MV
    - Functional dependencies can be used to check it
  - c) Aggregates must coincide with (or be computable from) the MV

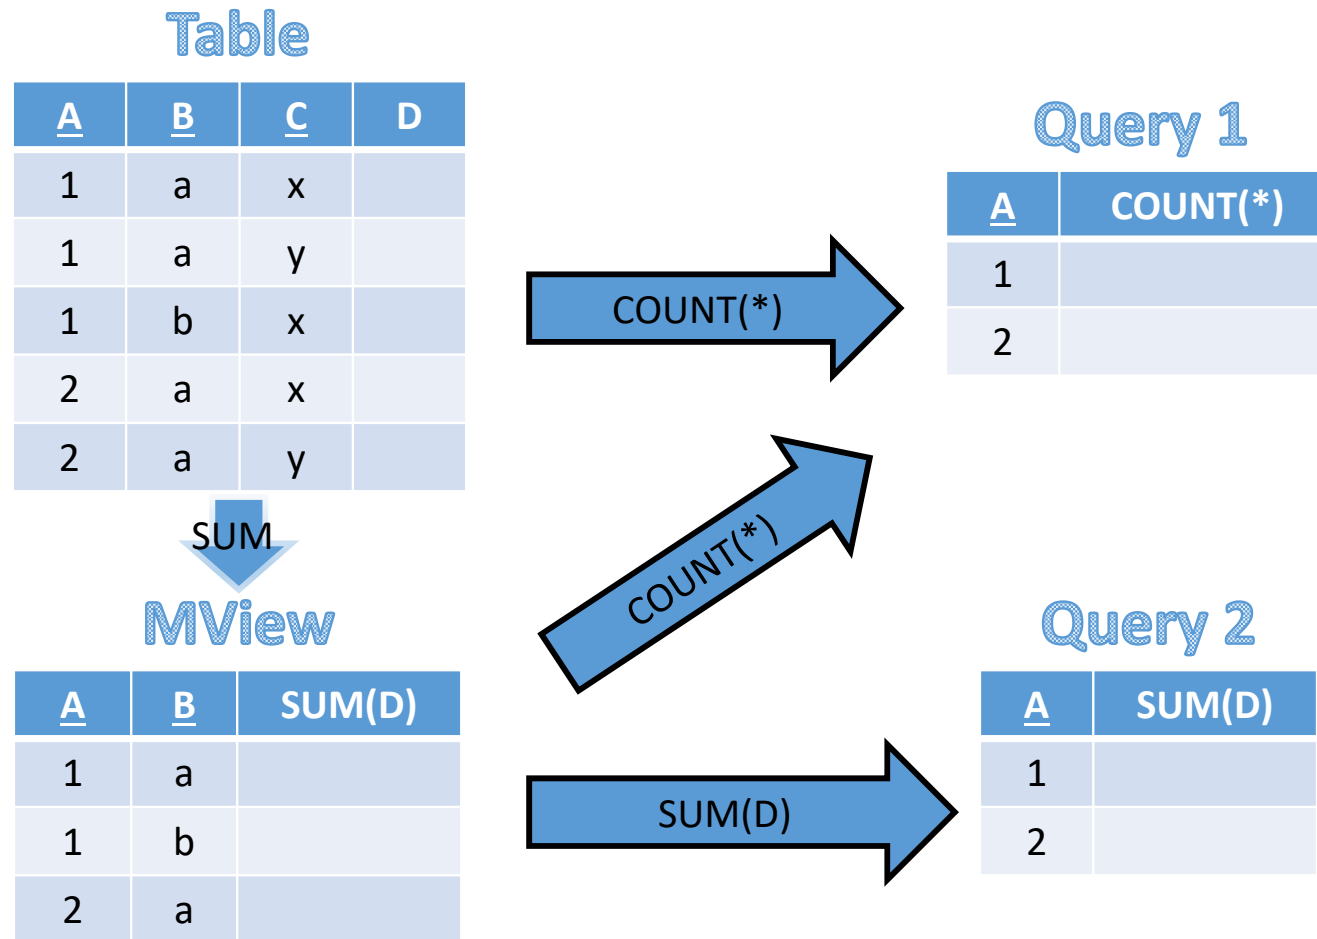
# Example of predicate subsumption requirement



# Example of aggregation level requirement



# Example of aggregation function requirement



# Materialized view selection

# To improve query performance ...


... Pre-compute as much as possible

- Redundant “tables” (a.k.a. materialized views)
  - Less attributes
  - Less tuples
    - Only those fulfilling the query predicate
    - Only one per combination of values of attributes in the GROUP BY
- Less space than the table
  - Less I/O to be accessed

# Problems in pre-computing

- Cost
  - Space
  - Time
    - Query vs Modification frequency
- Consistency and rewriting must be controlled
  - Using triggers
    - Advantages
      - Flexible
      - Allows rewriting of any query
      - Maybe efficient
    - Disadvantages
      - Complicates the management of the DBMS
      - Ad-hoc rewriting must be implemented for each query
        - Users are bound to our rewriting tools
  - Using MVs

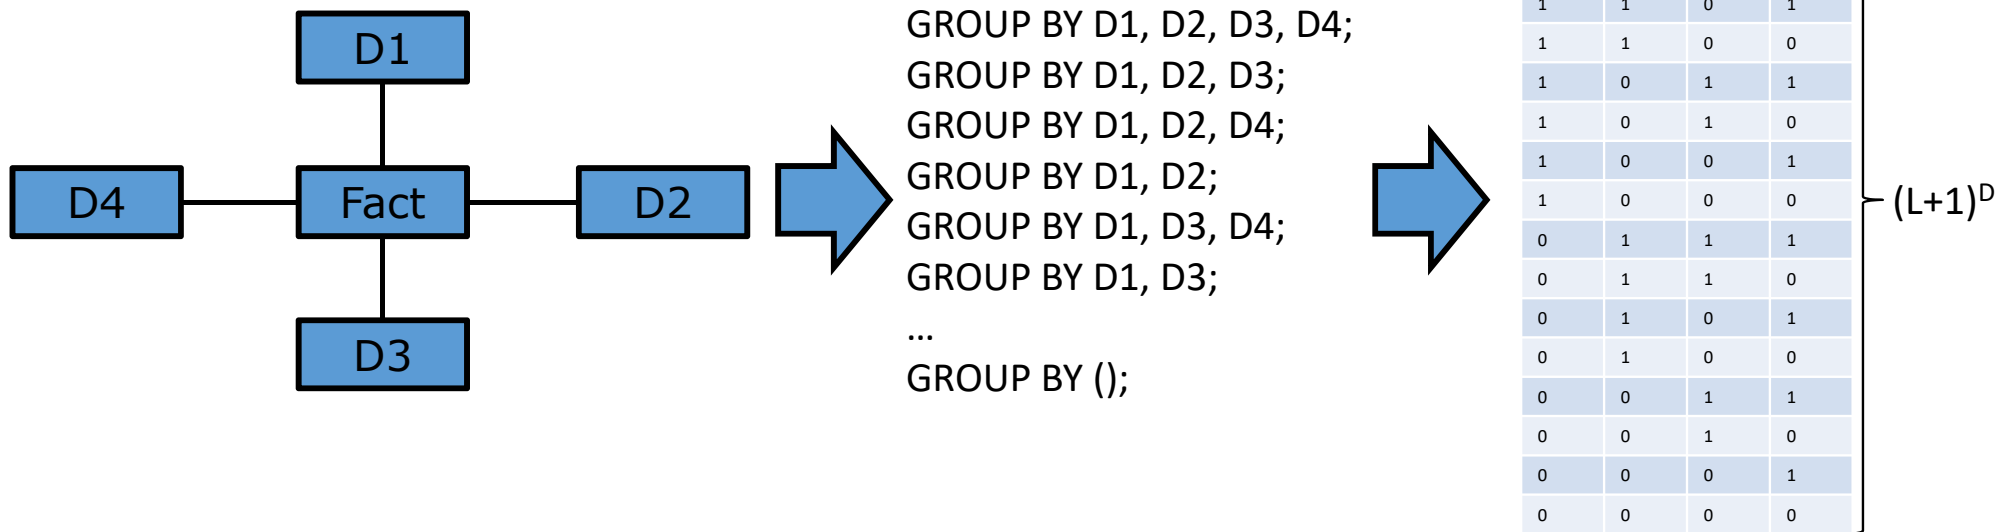
# Materialization trade-off






# Aggreion levels combinatorial explosion

- Choosing the best combination of views to be materialized is NP-complex
  - A fact table with  $D$  dimension tables with  $L$  aggregation levels (excluding the “All” level) for each one, would generate  $(L+1)^D$  possible MVs



# Solutions to the aggregation levels explosion

- The sparser the basic table/cube, (proportionally) the more space aggregates will use
  - Twelve days per year may generate twelve months per year
- Heuristics:
  - Materialize lower aggregation levels
    - They highly reduce the size (in absolute value) and solve many queries
  - Materialize higher aggregation levels
    - They are queried very often
  - Materialize a view if it solves a critical query or many queries
  - Do not materialize a view if it is a close successor of an already materialized view
    - Each tuple comes from the aggregation of at least 10
- Modify the set of MVs as user needs evolve



# Candidate views to be materialized

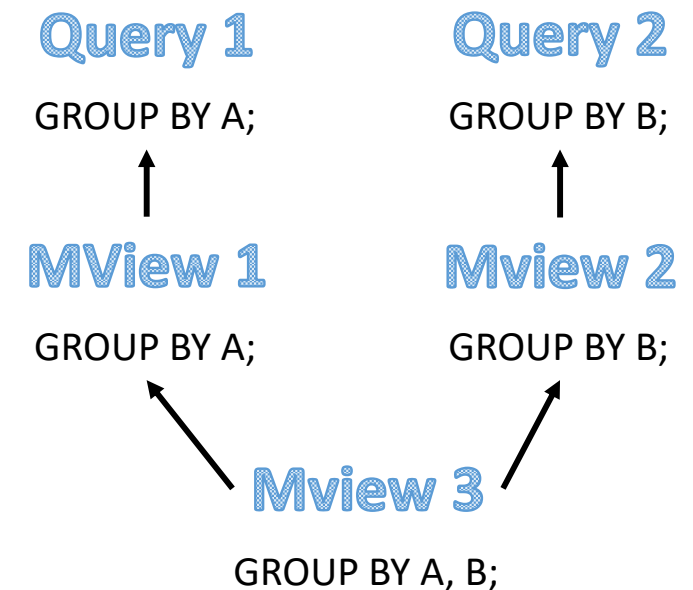
Given a workload  $W = \{q_1, q_2, q_3, \dots\}$ , and identifying queries by their GROUP BY clause, candidate views  $v_i$  are those that:

a)  $GB(v_i) = GB(q_j)$

b)  $GB(v_i) = \bigcup_{q_j \in Q} GB(q_j)$ , where  $Q \subseteq W$

Provided that:

- 1) Predicates also allow rewriting
- 2) Aggregations in the select clause also allow rewriting



# Algorithm to choose among candidates

Greedy algorithm (guarantees 63% minimum improvement):

Do

- 1) Consider those candidate views that fit in the available space and update time
- 2) Sort views based on the performance improvement they induce
- 3) Materialize first view in the list, if it improves performance

While performance improved and there is available space and update time

- Modify the set of MVs as user needs evolve

# Example of materialized view selection

Greedy

# Example of materialized view selection (I)

- Table CentMilResp(ref, pobl, edat, cand, val)
- $D=1\text{sec}$ ;  $C=0$
- $B_{\text{CentMilResp}}=10.000$ ;  $|\text{CentMilResp}|=100.000$
- $\text{Ndist}(\text{pobl})=200$ ;  $\text{Ndist}(\text{edat})=100$ ;  $\text{Ndist}(\text{cand})=10$
- All attributes require the same space
  - The control information (a.k.a. metadata) of the row requires as much space as another attribute
- Query frequencies:
  - 35%: `SELECT cand, MAX(val)`  
`FROM CentMilResp GROUP BY cand;`
  - 20%: `SELECT cand, edat, AVG(val), MAX(val), MIN(val)`  
`FROM CentMilResp GROUP BY cand, edat;`
  - 20%: `SELECT pobl, MAX(val)`  
`FROM CentMilResp GROUP BY cand, pobl;`
  - 25%: `SELECT pobl, MAX(val)`  
`FROM CentMilResp GROUP BY pobl;`
- We have 10.140 disk blocks available

# Example of materialized view selection (II)


**C1/Q1** - SELECT cand, MAX(val) FROM CentMilResp GROUP BY cand;

**C2/Q2** - SELECT cand, edat, AVG(val), MAX(val), MIN(val) FROM CentMilResp GROUP BY cand, edat;

**C3/Q3** - SELECT pobl, MAX(val) FROM CentMilResp GROUP BY cand, pobl;

**C4/Q4** - SELECT pobl, MAX(val) FROM CentMilResp GROUP BY pobl;

**C5** - SELECT cand, pobl, MAX(val) FROM CentMilResp GROUP BY cand, pobl;



Aggregation rows estimation:

$$|C_i| = \min(|T|, \text{Ndist}(a_1) * \dots * \text{Ndist}(a_n))$$

$$|C_2| = \min(100000, 10 * 100) = 1000$$

$$|C_3| = |C_5| = \min(100000, 10 * 200) = 2000$$

Materialized view space estimation:

$$B_{C_i} = \lceil B_T * (\text{Arity}(C_i) / \text{Arity}(T)) * (|C_i| / |T|) \rceil$$

$$B_{C_2} = \lceil 10000 * (6/6) * (1000/100000) \rceil = 100$$

$$B_{C_3} = \lceil 10000 * (3/6) * (2000/100000) \rceil = 100$$

$$B_{C_5} = \lceil 10000 * (4/6) * (2000/100000) \rceil = 134$$

$B_{\text{CentMilResp}}$	10000
$B_{C_1}$	1
$B_{C_2}$	100
$B_{C_3}$	100
$B_{C_4}$	10
$B_{C_5}$	134

# Example of materialized view selection (II)

**C1/Q1** - SELECT cand, MAX(val) FROM CentMilResp GROUP BY cand;

**C2/Q2** - SELECT cand, edat, AVG(val), MAX(val), MIN(val) FROM CentMilResp GROUP BY cand, edat;


**C3/Q3** - SELECT pobl, MAX(val) FROM CentMilResp GROUP BY cand, pobl;

**C4/Q4** - SELECT pobl, MAX(val) FROM CentMilResp GROUP BY pobl;

**C5** - SELECT cand, pobl, MAX(val) FROM CentMilResp GROUP BY cand, pobl;

Cost if there is no materialized view:

- Time: 10.000 sec/query
- Space: 10.000 blocks



	Q1 (35%)	Q2 (20%)	Q3 (20%)	Q4 (25%)	Avg
C1	1	10000	10000	10000	6500,4
C2	100	100	10000	10000	4555,0
C3	10000	10000	100	100	5545,0
C4	10000	10000	10000	10	7502,5
C5	134	10000	134	134	2107,2

$B_{\text{CentMilResp}}$	10000
$B_{C1}$	1
$B_{C2}$	100
$B_{C3}$	100
$B_{C4}$	10
$B_{C5}$	134



# Example of materialized view selection (III)


**C1/Q1** - SELECT cand, MAX(val) FROM CentMilResp GROUP BY cand;

**C2/Q2** - SELECT cand, edat, AVG(val), MAX(val), MIN(val) FROM CentMilResp GROUP BY cand, edat;

**C3/Q3** - SELECT pobl, MAX(val) FROM CentMilResp GROUP BY cand, pobl;

**C4/Q4** - SELECT pobl, MAX(val) FROM CentMilResp GROUP BY pobl;

**C5** - SELECT cand, pobl, MAX(val) FROM CentMilResp GROUP BY cand, pobl;



Cost if C5 is materialized:

- Time: 2.107,2 sec/query
- Space: 10.134 blocks

	Q1 (35%)	Q2 (20%)	Q3 (20%)	Q4 (25%)	Avg
C1	1	10000	134	134	2060,7
C2	100	100	134	134	115,3
C3	134	10000	100	100	2091,9
C4	134	10000	134	10	2076,2

Cost if C1 and C5 are materialized:

- Time: 2.060,7 sec/query
- Space: 10.135 blocks

B <sub>CentMilResp</sub>	10000
B <sub>C1</sub>	1
B <sub>C2</sub>	100
B <sub>C3</sub>	100
B <sub>C4</sub>	10
B <sub>C5</sub>	134

# Closing



# Summary

- ANSI/SPARC architecture
- Difficulties/Features when dealing with views
  - View expansion
  - Update through views
  - View updating
    - Assertions
  - Answering queries using views
    - Materialized view selection

# Bibliography

- J. Sistac et al. *Disseny de bases de dades*. Editorial UOC, 2002. Col·lecció Manuals, número 43
- G. Gardarin and P. Valduriez. *Relational databases and knowledge bases*. Addison Wesley Publishing Company, 1989
- T. Teorey et al. *Database modeling and design*, 4<sup>th</sup> edition. Morgan Kaufmann Publishers, 2006
- M. Golfarelli and S. Rizzi. *Data Warehouse Design*. McGraw-Hill, 2009